

Project 0

Warm-Up

INTRODUCTION:

It is not uncommon for students come into this course without basic C software development skills (which should have been developed in the CS111 prerequisites), and invest a great deal of time and frustrating effort in the first two projects before concluding they will not pass the course and must drop. We have created this simple warm-up to determine whether or not students are prepared to work on C programming projects. If you do not find this project to be relatively straight-forward, you may want to reconsider whether or not you are ready to take this course.

RELATION TO READING AND LECTURES:

None. This project requires only C programming skills and basic familiarity with command line interfaces and makefiles that incoming students should already possess. Mastery of the material covered in a class like 35L should be enough to make this project straightforward.

PROJECT OBJECTIVES:

- ensure students have a working GNU/Linux development environment.
- ensure students can code, compile, test and debug simple C programs.
- introduce and demonstrate the ability to use basic POSIX file operations.
- introduce and demonstrate the ability to process command line arguments.
- introduce and demonstrate the ability to catch and handle run-time exceptions.
- introduce and demonstrate the ability to return informative exit status.
- demonstrate the ability to research and exploit non-trivial APIs.
- demonstrate the ability to construct a standard `Makefile`.
- demonstrate the ability to write software that conforms to a Command Line Interface (CLI) specification.

DELIVERABLES:

A single compressed tarball (`.tar.gz`) containing:

- a single C source module that compiles cleanly (with `gcc`, with no errors or warnings).
- a `Makefile` to build the program and the tarball.
- two screen snapshot(s) from a `gdb(1)` session:
 - `backtrace.png` ... showing a segfault and associated stack-trace
 - `breakpoint.png` ... showing a breakpoint and variable inspection
- a `README` file with identification information, a description of the included files, the [smoke-test cases](#) in the `check` target, and any other information about your submission that you would like to bring to our attention (e.g., research, limitations, features, testing methodology).

PROJECT DESCRIPTION:

1. (if you do not already have one) bring up or obtain access to a GNU/Linux development environment. Your development environment should include (at least):
 - gcc
 - libc (e.g., glibc or libc6-dev)
 - make
 - gdb
2. (if you are not already familiar with them) study the following manual sections:
 - POSIX file operations ... *open(2)*, *creat(2)*, *close(2)*, *dup(2)*, *read(2)*, *write(2)*, *exit(2)*, *signal(2)*, and this brief tutorial on [file descriptor manipulation](#).
 - *strerror(3)* ... function that interprets the error codes returned from failed system calls.
 - *getopt(3)* ... the framework we will use for argument handling in all projects for this course.
 - *tar(1)* (and the *-z* option) ... program for archiving files in a tarball.
 - *gdb(1)* (and the *run*, *bt*, *list*, *print* and *break* commands in particular) ... a GNU/Linux debugger for C/C++ programs.

You will probably find understanding *getopt(3)* to be the most difficult part of this project. Feel free to seek out other examples/tutorials for these functions, but make sure you cite those sources in your `README`.

3. write a program that copies its standard input to its standard output by *read(2)*-ing from file descriptor 0 (until encountering an end of file) and *write(2)*-ing to file descriptor 1. If no errors (other than EOF) are encountered, your program should *exit(2)* with status 0.

Your program executable should be called `lab0`, and accept the following (optional) command line arguments (in any combination or order; duplicate uses of the same option with the same or different option-arguments are allowed):

- `--input=filename` ... use the specified file as standard input (instead of what is already being used as standard input)
If you are unable to open the specified input file, report the failure (on `stderr`, file descriptor 2) using *fprintf(3)*, and *exit(2)* with status 1.
- `--output=filename` ... create the specified file and use it as standard output (instead of whatever is already being used as standard output)
If the file already exists, truncate it to zero size. If you are unable to create or truncate the specified output file, report the failure (on `stderr`, file descriptor 2) using *fprintf(3)*, and *exit(2)* with status 2.
- `--segfault` ... force a segmentation fault (e.g., by calling a subroutine that sets a `char *` pointer to `NULL` and then stores through the null pointer). If this argument is specified, do it immediately, and do not copy from `stdin` to `stdout`.
- `--catch` ... use *signal(2)* to register a `SIGSEGV` handler that catches the segmentation fault, logs an error message (on `stderr`, file descriptor 2) and *exit(2)* with status 4.
- `--dump-core` ... dump core on segmentation faults.

When you print out an error message (e.g., because an open failed), your message should include enough information to enable a user to understand not merely the nature of the problem but its cause ... for example:

- which argument caused the problem ... e.g., `--input`
- which file could not be opened ... e.g., `myfile.txt`
- the reason it could not be opened ... e.g., `no such file`

Do your argument parsing with *getopt_long(3)*. This is, for historical reasons, a somewhat convoluted API, but ...

- it is very similar APIs are used in many other languages and systems.

- I want you to gain experience with the very common trial-and-error process of learning how to use a non-trivial API.

If you encounter an unrecognized argument you should print out an error message including a correct usage line, and *exit(2)* with status 3.

To ensure that operations are performed in the right order when multiple arguments are specified, it is suggested that you:

- First check all arguments for argument syntax.
- Then carry out all actions in the same order as the arguments.

It is relatively easy to generate primitive error messages with *perror(3)*, but if you study the documentation you will see how to get access to the underlying error descriptions, which you could then use with *fprintf(stderr, ...)* to generate better formatted error messages to the correct file descriptor.

Note that to use the advanced debugging features of *gdb(1)* you will need to compile your program with the *-g* option, which adds debugging symbol table information to your program.

4. create a *Makefile* that supports the following targets:
 - (default) ... build the *lab0* executable. To maximize compile-time error checking, you should compile your program with the *-Wall* and *-Wextra* options.
 - *check* ... runs a quick smoke-test on whether or not the program seems to work, supports the required arguments, and properly reports success or failure.
Please include a brief description (in your *README*) of what checks you chose to include in your smoke-test.
 - *clean* ... delete all files created by the *Makefile*, and return the directory to its freshly untared state.
 - *dist* ... build the distribution tarball.
5. run your program (with the *--segfault* argument) under *gdb(1)*
 - take the fault
 - get a stack backtrace
 - take a screen snapshot (to be included with your submission)

Putting the code that causes the *SEGFAULT* in a separate subroutine will make the stack trace a little more interesting.

6. run your program (with the *--segfault* argument) under *gdb(1)*
 - set a break-point at the bad assignment
 - run the program up to the breakpoint
 - inspect the pointer to confirm that it is indeed *NULL*
 - take a screen snapshot (to be included with your submission)

You will not be able to list lines of code, print data, or set breakpoints in *gdb* unless your program has been compiled with the *-g* (debug symbols) switch.

Summary of exit codes

- 0 ... copy successful
- 1 ... unrecognized argument
- 2 ... unable to open input file
- 3 ... unable to open output file
- 4 ... caught and received *SIGSEGV*
- 139 ... dumped core

SUBMISSION:

Your `README` file (for this and every lab) must include lines of the form:

- `NAME: your name`
- `EMAIL: your email`
- `ID: your student ID`

Your name, student ID, and email address should also appear as comments at the top of your `Makefile` and each source file.

Your tarball should have a name of the form `lab0-yourStudentID.tar.gz`. You can sanity check your submission with this [test script](#). There will be no manual regrading on this project. Submissions that do not pass the sanity check are likely to receive very low scores.

We will test it on a departmental GNU/Linux server. You would be well advised to test all the functionality of your submission on that platform before submitting it.

GRADING:

Points for this project will be awarded:

Value Feature

packaging and build

- 5% untars expected contents
- 5% clean build w/default action
- 1% correct `make check`
- 1% correct `make clean`
- 1% correct `make dist`
- 2% reasonability of `README` contents
- 5% reasonability of check smoke test

input/output features

- 10% correctly copy input to output
- 5% correctly implements `--input`
- 5% correctly implements `--output`
- 5% implements combined `--input + --output`
- 5% correct handling of un-openable/creatable output file
- 5% correct handling of non-existent input file
- 5% correct handling of invalid arguments

fault handling

- 5% generate (and die from) `SIGSEGV`
- 5% catch and report `SIGSEGV`

gdb use

- 5% screen shot showing taking of segfault within `gdb`
- 5% screen shot showing backtrace from segfault
- 5% screen shot showing breakpoint stop before fault

5% screen shot showing inspection of null pointer

code/package review

2% correct argument processing

2% correct file descriptor handling

2% correct signal handling

4% misc