

16. Virtual Memory

What methods of protecting memory from corruption by bad processes do we have?

1. Hire better programmers
2. Enable subscript/NULL checking
 - a. requires a language that stores array sizes
 - b. conditional branches and checking make this very slow
3. Base & Bound Registers
 - a. simple: only requires 2 registers
 - b. good for batch environments, bad for dynamic
 - c. requires contiguous memory per-process
 - i. external & internal fragmentation
 - d. not conducive to sharing
 - i. repetition of code with copying for each process
4. Segmented Memory
 - a. a hardware table mapping indices to a (array, index)/(segment#, offset) pair
 - b. does not require adjustment of locations
 - c. still has external fragmentation


All of these have downsides—a better and the most common one is virtual memory

This is for a few reasons:

1. frees programmers from worrying about physical locations & fragmentation
2. prevents processes from illegal memory access
3. lets processes and threads share memory
4. lets VM be larger than actual memory
 - a. while this is technically true, it can cause rapid switching of pages (**thrashing**)

How does it work?

- per thread register %cr3 points to the page table
- page table maps virtual to physical addresses
- Intel keeps the specific contents secret



Pages are supported by the hardware

- x86-64 has a CTRL register that stores page sizes
- How do we choose the size?
 - large pages → small table (saves time & space)
 - small pages → large table (less fragmentation)

Are loops allowed? (virtual→physical→virtual)

- this would allow us to look at the table
- Linux doesn't even let us look at the table to prevent memory cheating

How do we edit the page table then?

```
int mmap(void *vir_addr, size_t len, int prot, int flags, int
fd, off_t offset);
# returns the virtual address of the new page (not necessarily
the one you give)
# flags
  # fixed -- do not give me an address other than what I asked
for
  # private -- other processes are not allowed to see our memo
ry
```

What if the length isn't a multiple of the page size?

- ~ not allowed; use `int getpagesize()`
- ~ this limits pages to `INT_MAX`, so we use `long sysconf()`

Are all protections reasonable?

- `rw-`: ordinary data
- `r--`: read only data such as variables marked `const`
- `r-x`: code (read is given because they are open source)
- `rwX`: dynamically generated code
 - this allows buffer overflow attacks
 - the stack is often still `rwX` because of legacy
- `---`: useful for catching addressing errors

- can be used for guard pages on either side of the array
- can be used to guard a recursive frame from modifying the base pointer in stack
- IS used on 0 to protect from NULL accesses


~ x is expensive, so low end chips with no security risk is often combined with r

mmap trivia:

- the most common file to map is /dev/zero
 - this functions as an input only stream of zeroes
 - read ~ succeed & fill page with zeroes
 - write ~ succeed without doing anything
- suppose two processes try to map to the same address at the same
 - this is not allowed!
- if a file has been mapped, it will not be deleted while still mapped
- a mapped page is not intrinsically tied to any file

mmap allows us to use dynamic linking of libraries

```
$ gcc main.c -o executable -lm;
# -lm links to math linbrary
```



Processes share only an image of the code they share, so the virtual addresses do not match

— how do we map virtual addresses to physical addresses?

We use the Global Offset Table

- the GOT is rwx self-modifying code that gives the offset for machine language system calls