

21. Cloud Computing

This is a new field, so the terminology is not well defined — the ideas are what is important

We define cloud computing as :

- computation done on a network of servers maintained by a cloud service provider

With our model of cloud computing, our goal is effortless elasticity.


We can measure this with a few heuristics:

- scalability (ability to change size)
- availability (ability to run even with failure)
- fault tolerance (transparent availability)
- configuration (ease of setup)
- Quality of Service
- QoS Monitoring Quality

Cloud computing can be used to the benefit of IOT/Edge Computing.

- many home devices now connect to the internet
- devices which are now connected to the internet cannot do heavy computation

To speed up our computing we use Fog Computing



- localized nodes provide a hybrid between the device and the cloud
- the cloud functions like a cache for computation
- this forms a computation hierarchy

We want to make this transparent -- to hide the servers from developers entirely, so:

- we run provided functions in the servers
- developers are billed on usage, and it scales to \$0

Function at a Service (FaaS)

- a function is the primary unit of computation
- functions are executed in response to a trigger (such as an HTTP/HTTPS request)
- rules:
 - a function must execute relatively quickly
 - a function can have no persistent state
 - functions can, however, refer to storage or database servers to find a state)
- ex)


```
function main(params, context) {
  return { payload: 'Hello ' + params.name };
}

# params and context are JSON objects: a marshaling of a data
structure

# context is a JSON object that represents metadata like secur
ity
```

If we are worried about performance, we would never use FaaS: (un)marshaling is slow

- we use FaaS for processing speed or functionality



Note that the workers are virtual machines implemented by Docker (for example)

- the developers believe the STATE is not part of the QoS metrics — the providers disagree

The steps of event execution :

1. arrival
2. validation
 - a. authentication
 - b. authorization
 - c. resource limit checking
3. enqueue

4. dispatch
5. allocate a container (a cheap VM)
6. copy function code into the container *
7. execute the function
8. deallocate the container

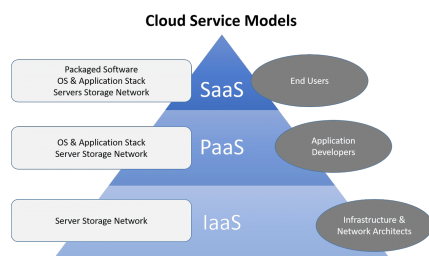
* our bottleneck is step (6) since Linux booting is slow relative to all the other operations

- this is called the cold start problem;
 - this causes our latency to skyrocket
 - we can temper the throughput by adding more workers
 - How do we temper this?
 - pre-load stem cell containers with modules we expect to use
 - do not clear warm containers post use
 - We can get to an almost reasonable speed with these techniques!
- we call this model the Action and Trigger Model
 - a model where an event triggers the execution of an action
 - one event can trigger multiple actions
 - we can use this to implement parallel execution
 - one action can cause an event that triggers an action
 - we can use this to implement serialized execution
- this has problems:
 - debugging
 - GDB is too big to use! We have to use their logs!
 - fixing bottlenecks
 - We can't use ps! We have to check usage in the logs post-run!
 - atomicity(?)
 - We are sometimes guaranteed atomicity of function calls
 - change
 - We have difficulty refactoring (breaking functions) and reverting to old versions
 - The tools are evolving too quickly for us to get used to them!

We have a few basic models of cloud computing which provide different levels of services

- Infrastructure as a Service (IaaS)

- the provides supplies virtual machines
 - a stripped down OS called a hypervisor hosts each VM
 - common hypervisors include Zen and Virtual Box
 - the programmer is responsible for the OS, configuration, and applications
 - at the enterprise level, this is 2/3 of IT spending
 - it will not reach 100% because of security concerns
 - there are some major issues, though:
 - we have to anticipate resource usage
 - it is hard to adjust usage on the fly
- Platform as a Service (PaaS)
 - the provider supplies some of the software stack in addition to VMs
 - less expertise is needed, but provisioning is still pushed on the developer
 - developers still have to autoscale
- Backend as a Service (BaaS)
 - the provider supplies a framework for an (often mobile) app on the backend and client
 - avoids heavy computation on mobile devices
 - scales, but allows little flexibility
- Software as a Service (SaaS)
 - provider supplies the entire platform & stack
 - the developer calls the service's API (which is application specific)
 - can be tailored with callbacks in config files



Our models form a pyramid of increasing ease but decreasing flexibility