

Why do we use signals?

- asynchronous I/O with `aread()`
  - returns right away and kernel keeps going with handling the data
  - get a `SIGIO` signal later
- error in your code (`SIGFPE`)
  - divide by zero
  - floating point overflow
  - invalid instruction
- impatient user of infinite loop (`SIGINT`)
  - `^C` to end program
- impending power outage (`SIGPWR`)
  - So we can do any saving before shutdown
- to check for dying children (`SIGCHLD`)
  - `p = waitpid(-1, &status, WNOHANG)`
  - now we don't have to call this method every 100 milliseconds
- user went away (`SIGHUP`)
- alarms
  - alarms are not inherited by `fork()` but by `execvp()`
- suspending a process
  - `$ kill -STOP 29; kill -CONT 29`
- kill the program `SIGKILL`
  - cannot be caught or ignored

We can kill a process with:

```
while (fork()) continue;
$ kill -KILL 29316      # does not kill children
# however this does not kill the shell bomb
$ kill -KILL -29316     # kills all children as well
```

Code must often be developed specifically to be able to handle signals; take the code

```
fd = open("foo", O_RDONLY);
fo = open("foo.gz", O_WRONLY);
while (compress(fd, fo)) continue;
close(fo);
unlink("foo");
// THIS CODE IS NOT ATOMIC AND CAN BE INTERRUPTED BY A SIGNAL
```

We can attempt to avoid these errors by implementing a signal handler

```
static void cleanup (int sig) {
    unlink("foo.gz");
    _exit(1);
}
int main() { ...
    fd = open("foo", O_RDONLY);
    signal(SIGINT, cleanup);*
    fo = open("foo.gz", O_WRONLY);
    while (compress(fd, fo)) continue;
    close(fo);
    signal(SIGINT, SIG_DFL);*
    unlink("foo");
    ... }
```

but this still leaves race conditions .

In our current implementation, all threads are affected by the sign. Should all threads handle the signal? Would all threads handle it the same? NO; threads have their own signal mask to ignore signals. This is why `pthread_sigmask()` affects only current thread!

So how do we handle them? By default threads have their signals blocked. We use `pthread_sigmask()` to unblock the signal if we want a thread to handle it. Linux picks one random thread to deliver the signal. We make a mask with

```
int pthread_sigmask(int how, const sigset_t *set, sigset_t *oldset);
// how = SIG_BLOCK, SIG_UNBLOCK, SIG_SETMASK
```

this allows the signal to arrive even before function returns. With this, we can build critical sections such as

```
sigset_t ss;
sigemptyset(&ss);
sigaddset(&ss, SIGINT);
pthread_sigmask(SIG_BLOCK, &ss, 0);
// critical section here .....
pthread_sigmask(SIG_UNBLOCK, &ss, 0);
```

But how can a signal handler manage memory access?

```
void handle_interrupt(int sig) {
    fprintf(stderr, "Interrupted\n");
    unlink(...);
}
fprintf(...) { malloc(...); } // interrupt
malloc(...) { // operating on heap }
// if we interrupt malloc and fprintf will call malloc again
// the second malloc may corrupt the heap, thus the first malloc call
```

Only some system calls can safely be used in handlers! We can call most system calls, such as:

- `_exit()`
- `write()`

But there are exceptions:

- `exit()` (calls `malloc`, flushes I/O buffer)
- `fprintf()`
- `malloc()`

We can perform all system calls in a single handler with:

```
void handle_interrupt(int sig) {
    if (pthread_self() == stgmgr) really_handle_interrupt();
    else pthread_kill(SIGINT, stgmgr); // forward signal to stage manager
}
```

```
# a more conservative approach is to set the variable and handle outside
sig_atomic_t volatile globv;
void handle_interrupt(int sig) {
    global = 1;
}
// always memory access, no cache!
```

But even with our scrupulous effort, interrupts can still cause difficulty:

```
read("/dev/tty", buf, 100);  
// SIGHUP signal arrives  
// run SIGHUP handler  
// returns and continue reading
```

This means we have to complicate our code:

```
while (read("/dev/tty", buf, 100) == -1 && errno == EINTR) continue;
```

These types of errors are common with long system calls; clearly scheduling concurrent threads properly is important.