

Schedulers can work on multiple levels:

1. long-term scheduler
which processes should be admitted to OS?
 - look at return value of `fork()`
 - if `fork() < 0`, denied by OS
2. medium-term scheduler
which processes reside in RAM?
3. short-term scheduler
which threads get to run on the limited number of CPUs

A process can be in one of three states:

1. *Running* := executing instructions on a processor
2. *Ready* := ready to run but waiting on OS
3. *Blocked* := waiting on another event

In addition to these three, there are the edge cases of:

1. *Initial* := just created, environment not set up
2. *Final/Zombie* := process is complete but hasn't been cleaned up yet

Moving from ready to running is called being *scheduled*. Moving from running to ready is called being *de-scheduled*. The act of scheduling is assigning CPUs to threads. Each instruction pointer requires a CPU to run. This requires a harmony of hardware and software

Determining how to schedule is easy when we have enough CPUs, but what do we do if there are more threads than CPUs?

We need a few things to answer this

- Theory: scheduling policy
- Practice: Scheduling and Dispatch Mechanisms

Tiny gaps where neither thread is executing are called *context switches*, and are done by the OS. These can be timed in two broad ways:

1. Cooperative Scheduling := a thread “volunteers” to give up its CPU with syscalls
2. Preemptive Scheduling := the OS preempts threads every time slice
 - this is cheap and easier on the OS
 - this is common in IOT/embedded systems
 - short slice = less efficient
 - long slice = long wait

How do threads “volunteer”?

```
#include <sched.h>
int sched_yield(void);
```

We can use this to make waiting on a device more efficient:

```
// we can bust wait
while(isbusy(device)) continue;
// we can poll, which is still ineffecient (but better)
while(isbusy(device)) sched_yield()
// or we can block, telling kernel to wake it when the condition is met
while(isbusy(device)) wait_until_ready()
```


The Linux scheduler runs the following algorithm:

```

for (;;) {
    choose an unblocked thread
    load it into a CPU
    run it
    yield
    store state
    for each thread that has become ready
        unblock
}

```

This scheduler is completely terrible! People have measured! But what did they measure?



In addition, SEASNET screws us students over with a priority queue with priorities:

1. root
2. operations staff
3. students/faculty

And each of these contains its own sub-scheduler and algorithm

Usually, priority 1 is the highest priority, but SEASNET uses the idea of *niceness*. If p1 has niceness x and p2 has niceness $y > x$, p2 will defer. users can raise a program's niceness, but cannot lower it.

We can't be too harsh though... there is a lot of complexity involved in *real-time systems*.

0.1 Real-Time Scheduling

This contains two types of deadlines:

- HARD:
 - deadlines CANNOT be missed, \implies performance = correctness
 - predictability $>$ performance, \implies caches are the enemy
 - use polling instead of interrupts, since polling controls test duration
- SOFT:
 - a missed deadline is not necessarily a failure
 - 2 scheduling options:
 1. rate-monotonic scheduling := give a % usage to job data streams
 2. earliest deadline first := can drop late requests when inundated. This allows one stream to monopolize the CPU.
 - ex) Video Playback: Each frame is treated as its own request. When the connection is slow, the scheduler periodically drops frames; this is often imperceptible to humans

Most real-life systems have both hard and soft real-time scheduling; this introduces a lot of complexity.